

## Here's what messy data look like

---

```
# View the first 6 rows of data
```

```
head(weather, n=6)
```

```
# View the last 6 rows of data
```

```
tail(weather , n=6)
```

```
# View a condensed summary of the data
```

```
str(weather)
```

## Here's what clean data look like

---

```
# View the first 6 rows of data
```

```
head(weather_clean ,n=6)
```

```
# View the last 6 rows of data
```

```
tail(weather_clean , n=6)
```

```
# View a condensed summary of the data
```

```
str(weather_clean)
```

## Getting a feel for your data

---

```
# Check the class of bmi
```

```
class(bmi)
```

```
# Check the dimensions of bmi
```

```
dim(bmi)
```

```
# View the column names of bmi
```

```
names(bmi)
```

## Viewing the structure of your data

---

```
# Check the structure of bmi
```

```
str(bmi)
```

```
# Load dplyr
```

```
library(dplyr)
```

```
# Check the structure of bmi, the dplyr way
```

```
glimpse(bmi)
```

```
# View a summary of bmi
```

```
summary(bmi)
```

## Looking at your data

---

```
# Print bmi to the console
```

```
print(bmi)
```

```
# View the first 6 rows
```

```
head(bmi)
```

```
# View the first 15 rows
```

```
head(bmi,n=15)
```

```
# View the last 6 rows
```

```
tail(bmi)
```

```
# View the last 10 rows
```

```
tail(bmi,n=10)
```

## Visualizing your data

---

```
# Histogram of BMIs from 2008
```

```
hist(bmi$Y2008)
```

```
# Scatter plot comparing BMIs from 1980 to those from 2008
```

```
plot(bmi$Y1980,bmi$Y2008)
```

## Gathering columns into key-value pairs

---

```
# Apply gather() to bmi and save the result as bmi_long
```

```
bmi_long <- gather(bmi, year, value, -Country)
```

```
# View the first 20 rows of the result
```

```
head(bmi_long,n=20)
```

## Spreading key-value pairs into columns

---

```
# Apply spread() to bmi_long
```

```
bmi_wide <- spread(bmi_long, year, bmi_val)
```

```
# View the head of bmi_wide
```

```
head(bmi_wide)
```

## Separating columns

---

```
# Apply separate() to bmi_cc
bmi_cc_clean <- separate(bmi_cc, col = Country_ISO, into = c("Country", "ISO"), sep = "/")

# Print the head of the result
head(bmi_cc_clean)
```

## Uniting columns

---

```
# Apply unite() to bmi_cc_clean
bmi_cc <- unite(bmi_cc_clean, Country_ISO, Country, ISO, sep = "-")

# View the head of the result
head(bmi_cc)
```

## Column headers are values, not variable names

---

```
## tidyr and dplyr are already loaded for you

# View the head of census
head(census)

# Gather the month columns
census2 <- gather(census, month, amount, -YEAR)

# Arrange rows by YEAR using dplyr's arrange
census2 <- arrange(census2, YEAR)
```

```
# View first 20 rows of census2
```

```
head(census2,n=20)
```

## Variables are stored in both rows and columns

---

```
## tidyr is already loaded for you
```

```
# View first 50 rows of census_long
```

```
head(census_long,n=50)
```

```
# Spread the type column
```

```
census_long2 <- spread(census_long, type, amount)
```

```
# View first 20 rows of census_long2
```

```
head(census_long2,n=20)
```

## Multiple values are stored in one column

---

```
## tidyr is already loaded for you
```

```
# View the head of census_long3
```

```
head(census_long3)
```

```
# Separate the yr_month column into two
```

```
census_long4 <- separate(census_long3, yr_month, c("year", "month"))
```

```
# View the first 6 rows of the result
```

```
head(census_long4)
```

## Types of variables in R

---

```
# Make this evaluate to character
```

```
class("true")
```

```
a <- as.character("true")
```

```
# Make this evaluate to numeric
```

```
class(8484.00)
```

```
b <- as.numeric(8484.00)
```

```
# Make this evaluate to integer
```

```
class(99L)
```

```
c <- as.integer(99)
```

```
# Make this evaluate to factor
```

```
class(factor("factor"))
```

```
d <- as.factor("factor")
```

```
# Make this evaluate to logical
```

```
class(FALSE)
```

```
e <- as.logical(FALSE)
```

## Working with dates

---

```
# Preview students2 with str()
```

```
str(students2)
```

```
# Load the lubridate package
library(lubridate)

# Parse as date
dmy("17 Sep 2015")

# Parse as date and time (with no seconds!)
mdy_hm("July 15, 2012 12:56")

# Coerce dob to a date (with no time)
students2$dob <- ymd(students2$dob)

# Coerce nurse_visit to a date and time
students2$nurse_visit <- ymd_hms(students2$nurse_visit)

# Look at students2 once more with str()
str(students2)
```

## Trimming and padding strings

---

```
# Load the stringr package
library(stringr)

# Trim all leading and trailing whitespace
str_trim(c(" Filip ", "Nick ", " Jonathan"))

# Pad these strings with leading zeros
str_pad(c("23485W", "8823453Q", "994Z"),width = 9, side = "left", pad = "0")
```

## Upper and lower case

---

```
# Print state abbreviations
print(states)

# Make states all uppercase and save result to states_upper
states_upper <- toupper(states)

# Make states_upper all lowercase again
tolower(states)
```

## Finding and replacing strings

---

```
## stringr has been loaded for you

# Look at the head of students2
head(students2)

# Detect all dates of birth (dob) in 1997
str_detect(students2$dob, "1997")

# In the sex column, replace "F" with "Female"...
students2$sex <- str_replace(students2$sex, "F", "Female")

# ...And "M" with "Male"
students2$sex <- str_replace(students2$sex, "M", "Male")

# View the head of students2
head(students2)
```



## Finding missing values

---

# Call is.na() on the full social\_df to spot all NAs

```
is.na(social_df)
```

# Use the any() function to ask whether there are any NAs in the data

```
any(is.na(social_df))
```

# View a summary() of the dataset

```
summary(social_df)
```

# Call table() on the status column

```
table(social_df$status)
```

## Dealing with missing values

---

## The stringr package is preloaded

# Replace all empty strings in status with NA

```
social_df$status[social_df$status == ""] <- NA
```

# Print social\_df to the console

```
social_df
```

# Use complete.cases() to see which rows have no missing values

```
complete.cases(social_df)
```

```
# Use na.omit() to remove all rows with any missing values
na.omit(social_df)
```

## **Outliers and obvious errors**

### **Dealing with outliers and obvious errors**

---

```
# Look at a summary() of students3
summary(students3)
```

```
# View a histogram of the age variable
hist(students3$age)
```

```
# View a histogram of the absences variable
hist(students3$absences)
```

```
# View a histogram of absences, but force zeros to be bucketed to the right of zero
hist(students3$absences, right = FALSE)
```

```
# View a boxplot of age
boxplot(students3$age)
```

```
# View a boxplot of absences
boxplot(students3$absences)
```

## **Case Study: Get a feel for the data**

---

```
# Verify that weather is a data.frame
class(weather)
```

```
# Check the dimensions
```

```
dim(weather)
```

```
# View the column names
```

```
names(weather)
```

## Summarize the data

---

```
# View the structure of the data
```

```
str(weather)
```

```
# Load dplyr package
```

```
library(dplyr)
```

```
# Look at the structure using dplyr's glimpse()
```

```
glimpse(weather)
```

```
# View a summary of the data
```

```
summary(weather)
```

## Take a closer look

### Column names are values

---

```
# Load the tidyr package
```

```
library(tidyr)
```

```
# Gather the columns
```

```
weather2 <- gather(weather, day, value, X1:X31, -year, na.rm = TRUE)
```

```
# View the head
```

```
head(weather2)
```

## Values are variable names

---

```
## The tidyr package is already loaded
```

```
# First remove column of row names
```

```
weather2 <- weather2[, -1]
```

```
# Spread the data
```

```
weather3 <- spread(weather2, measure, value)
```

```
# View the head
```

```
head(weather3)
```

## Clean up dates

---

```
## tidyr and dplyr are already loaded
```

```
# Load the stringr and lubridate packages
```

```
library(stringr)
```

```
library(lubridate)
```

```
# Remove X's from day column
```

```
weather3$day <- str_replace(weather3$day, "X", "")
```

```
# Unite the year, month, and day columns
weather4 <- unite(weather3, date, year, month, day, sep = "-")

# Convert date column to proper date format using lubridates's ymd()
weather4$date <- ymd(weather4$date)

# Rearrange columns using dplyr's select()
weather5 <- select(weather4, date, Events, CloudCover:WindDirDegrees)

# View the head of weather5
head(weather5)
```

## A closer look at column types

---

```
# View the structure of weather5
str(weather5)

# Examine the first 20 rows of weather5. Are most of the characters numeric?
head(weather5, n=20)

# See what happens if we try to convert PrecipitationIn to numeric
as.numeric(weather5$PrecipitationIn)
```

## Column type conversions

---

```
## The dplyr and stringr packages are already loaded

# Replace T with 0 (T = trace)
weather5$PrecipitationIn <- str_replace(weather5$PrecipitationIn, "T", 0)
```

```
# Convert characters to numerics
```

```
weather6 <- mutate_each(weather5, funs(as.numeric), CloudCover:WindDirDegrees)
```

```
# Look at result
```

```
str(weather6)
```

## Find missing values

---

```
# Count missing values
```

```
sum(is.na(weather6))
```

```
# Find missing values
```

```
summary(weather6)
```

```
# Find indices of NAs in Max.Gust.SpeedMPH
```

```
ind <- which(is.na(weather6$Max.Gust.SpeedMPH))
```

```
# Look at the full rows for records missing Max.Gust.SpeedMPH
```

```
weather6[ind, ]
```

## An obvious error

---

```
# Review distributions for all variables
```

```
summary(weather6)
```

```
# Find row with Max.Humidity of 1000
```

```
ind <- which(weather6$Max.Humidity == 1000)
```

```
# Look at the data for that day
```

```
weather6[ind, ]
```

```
# Change 1000 to 100
```

```
weather6$Max.Humidity[ind] <- 100
```

## Another obvious error

---

```
# Look at summary of Mean.VisibilityMiles
```

```
summary(weather6$Mean.VisibilityMiles)
```

```
# Get index of row with -1 value
```

```
ind <- which(weather6$Mean.VisibilityMiles == -1)
```

```
# Look at full row
```

```
weather6[ind, ]
```

```
# Set Mean.VisibilityMiles to the appropriate value
```

```
weather6$Mean.VisibilityMiles[ind] <- 10
```

## Check other extreme values

---

```
# Review summary of full data once more
```

```
summary(weather6 )
```

```
# Look at histogram for MeanDew.PointF
```

```
hist(weather6$MeanDew.PointF)
```

```
# Look at histogram for Min.TemperatureF
```

```
hist(weather6$Min.TemperatureF)
```

```
# Compare to histogram for Mean.TemperatureF
```

```
hist(weather6$Mean.TemperatureF)
```

## Finishing touches

---

```
# Clean up column names
```

```
names(weather6) <- new_colnames
```

```
# Replace empty cells in events column
```

```
weather6$events[weather6$events == ""] <- "None"
```

```
# Print the first 6 rows of weather6
```

```
head(weather6)
```